



Deep neural networks and back-propagation


王兴刚

<https://xinggangw.info>

Neural network



$$g(\mathbf{x}) = f\left(\sum_{i=1}^d x_i w_i + w_0\right) = f(\mathbf{w}^t \mathbf{x})$$




$f(net)$

Activation functions


Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$




tanh

$$\tanh(x)$$




ReLU

$$\max(0, x)$$



Leaky ReLU

$$\max(0.1x, x)$$




Maxout


$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

ELU

$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$




Neural networks: Architectures



“2-layer Neural Net”, or
“1-hidden-layer Neural Net”


“Fully-connected” layers



“3-layer Neural Net”, or
“2-hidden-layer Neural Net”

Forward

$$\begin{aligned}z_1 &= x^T W_1 + b_1 \\a_1 &= \tanh(z_1) \\z_2 &= a_1^T W_2 + b_2 \\ \hat{y} &= a_2 = \text{softmax}(z_2)\end{aligned}$$



$$z_2 = [s_1, s_2, \dots, s_C]$$

$$\text{softmax}(s_k | z_2) = \frac{e^{s_k}}{\sum_{j=1}^C e^{s_j}}$$

Loss function



$$L_i = -\log\left(\frac{e^{s_{y_i}}}{\sum_j e^{s_j}}\right)$$

unnormalized probabilities

cat
car
frog

3.2
5.1
-1.7

exp

24.5
164.0
0.18

normalize

0.13
0.87
0.00

$$L_i = -\log(0.13) \\ = 0.89$$

unnormalized log probabilities

probabilities

$$L(y, \hat{y}) = -\frac{1}{N} \sum_{n \in N} \sum_{i \in C} y_{n,i} \log \hat{y}_{n,i}$$

Back propagation

Then the derivation of the softmax cross entropy loss is given as follows.

Let o_k denote the k -th node of the input layer of the following softmax layer.
The calculation of softmax function is given as follows.

$$p_j = \frac{e^{o_j}}{\sum_k e^{o_k}} \quad (1)$$

The standard cross entropy loss function L is given as follows.

$$L = - \sum_j y_j \log p_j, \quad (2)$$

$$\frac{\partial L}{\partial o_i} = - \sum_k y_k \frac{\partial \log p_k}{\partial o_i} \quad (3)$$

$$= - \sum_k y_k \frac{1}{p_k} \frac{\partial p_k}{\partial o_i} \quad (4)$$

$$= -y_i(1 - p_i) - \sum_{k \neq i} y_k \frac{1}{p_k} (-p_k p_i) \quad (5)$$

$$= -y_i(1 - p_i) + \sum_{k \neq i} y_k (p_i) \quad (6)$$

$$= -y_i + y_i p_i + \sum_{k \neq i} y_k (p_i) \quad (7)$$

$$= p_i \left(\sum_k y_k \right) - y_i \quad (8)$$

$$= p_i - y_i \quad (9)$$

Back propagation

Forward:

$$z_1 = x^T W_1 + b_1$$

$$a_1 = \tanh(z_1)$$

$$z_2 = a_1^T W_2 + b_2$$

$$\hat{y} = a_2 = \text{softmax}(z_2)$$

Backward:

$$\delta_3 = \hat{y} - y$$

$$\delta_2 = (1 - \tanh^2 z_1) \circ \delta_3 W_2^T$$

$$\frac{\partial L}{\partial W_2} = a_1^T \delta_3$$

$$\frac{\partial L}{\partial b_2} = \delta_3$$

$$\frac{\partial L}{\partial W_1} = x^T \delta_2$$

$$\frac{\partial L}{\partial b_1} = \delta_2$$

Back propagation

Forward:

$$z_1 = x^T W_1 + b_1$$

$$a_1 = \tanh(z_1)$$

$$z_2 = a_1^T W_2 + b_2$$

$$\hat{y} = a_2 = \text{softmax}(z_2)$$

Backward:

$$\frac{\partial L}{\partial W_2} = \frac{\partial L}{\partial z_2} \frac{\partial z_2}{\partial W_2} = (\hat{y} - y) a_1^T$$

$$\frac{\partial L}{\partial b_2} = \frac{\partial L}{\partial z_2} \frac{\partial z_2}{\partial b_2} = (\hat{y} - y)$$

$$\frac{\partial L}{\partial W_1} = \frac{\partial L}{\partial z_2} \frac{\partial z_2}{\partial a_1} \frac{\partial a_1}{\partial z_1} \frac{\partial z_1}{\partial W_1}$$

$(\hat{y} - y) \quad W_2 \quad (1 - \tanh^2 z_1) \quad x^T$

Back propagation

Forward:

$$z_1 = x^T W_1 + b_1$$

$$a_1 = \tanh(z_1)$$

$$z_2 = a_1^T W_2 + b_2$$

$$\hat{y} = a_2 = \text{softmax}(z_2)$$

Backward:

$$\delta_3 = \hat{y} - y$$

$$\delta_2 = (1 - \tanh^2 z_1) \circ \delta_3 W_2^T$$

$$\frac{\partial L}{\partial W_2} = a_1^T \delta_3$$

$$\frac{\partial L}{\partial b_2} = \delta_3$$

$$\frac{\partial L}{\partial W_1} = x^T \delta_2$$

$$\frac{\partial L}{\partial b_1} = \delta_2$$

Implementation

Forward:

```
1 # Helper function to predict an output (0 or 1)
2 def predict(model, x):
3     W1, b1, W2, b2 = model['W1'], model['b1'], model['W2'], model['b2']
4     # Forward propagation
5     z1 = x.dot(W1) + b1
6     a1 = np.tanh(z1)
7     z2 = a1.dot(W2) + b2
8     exp_scores = np.exp(z2)
9     probs = exp_scores / np.sum(exp_scores, axis=1, keepdims=True)
10    return np.argmax(probs, axis=1)
```

Implementation


Backward:

```
17 # Gradient descent. For each batch...
18 for i in xrange(0, num_passes):
19
20     # Forward propagation
21     z1 = X.dot(W1) + b1
22     a1 = np.tanh(z1)
23     z2 = a1.dot(W2) + b2
24     exp_scores = np.exp(z2)
25     probs = exp_scores / np.sum(exp_scores, axis=1, keepdims=True)
26
27     # Backpropagation
28     delta3 = probs
29     delta3[range(num_examples), y] -= 1
30     dW2 = (a1.T).dot(delta3)
31     db2 = np.sum(delta3, axis=0, keepdims=True)
32     delta2 = delta3.dot(W2.T) * (1 - np.power(a1, 2))
33     dW1 = np.dot(X.T, delta2)
34     db1 = np.sum(delta2, axis=0)
35
36     # Add regularization terms (b1 and b2 don't have regularization terms)
37     dW2 += reg_lambda * W2
38     dW1 += reg_lambda * W1
39
40     # Gradient descent parameter update
41     W1 += -epsilon * dW1
42     b1 += -epsilon * db1
43     W2 += -epsilon * dW2
44     b2 += -epsilon * db2
```


```
while True:
    data_batch = dataset.sample_data_batch()
    loss = network.forward(data_batch)
    dx = network.backward()
    x += - learning_rate * dx
```

Results

Logistic Regression



Decision Boundary for hidden layer size 5



L0 regularization

$$\text{Cost} = \sum_{i=0}^N (y_i - \sum_{j=0}^M x_{ij} \mathbf{w}_j)^2 + \lambda ||\mathbf{w}||_0$$

$||\mathbf{w}||_0$ means non-zero elements in w

$||\mathbf{w}||_0$ is non-convex and not differentiable.

L1 regularization

$$L'(\mathbf{w}) = \sum_{i=0}^N (y_i - \sum_{j=0}^M x_{ij} \mathbf{w}_j)^2 + \lambda |\mathbf{w}|_1$$

L1 regularization: $|\mathbf{w}|_1 = |w_1| + |w_2| + \dots$

$$L'(\mathbf{w}) = L(\mathbf{w}) + \lambda |\mathbf{w}|_1 \quad \frac{\partial L'}{\partial \mathbf{w}} = \frac{\partial L}{\partial \mathbf{w}} + \lambda \text{sign}(\mathbf{w})$$

Parameter update:

$$\begin{aligned} \mathbf{w}^{t+1} &= \mathbf{w}^t - \eta \frac{\partial L'}{\partial \mathbf{w}} = \mathbf{w}^t - \eta \left(\frac{\partial L}{\partial \mathbf{w}} + \lambda \text{sign}(\mathbf{w}) \right) \\ &= \mathbf{w}^t - \eta \frac{\partial L}{\partial \mathbf{w}} - \eta \lambda \text{sign}(\mathbf{w}) \end{aligned}$$

$\eta \lambda \text{sign}(\mathbf{w})$ always makes weight smaller (closing to zero)

L2 regularization

$$L'(\mathbf{w}) = \sum_{i=0}^N (y_i - \sum_{j=0}^M x_{ij} \mathbf{w}_j)^2 + \lambda \frac{1}{2} |\mathbf{w}|_2$$

L1 regularization: $|\mathbf{w}|_2 = |w_1|^2 + |w_2|^2 + \dots$

$$L'(\mathbf{w}) = L(\mathbf{w}) + \lambda |\mathbf{w}|_2 \qquad \frac{\partial L'}{\partial \mathbf{w}} = \frac{\partial L}{\partial \mathbf{w}} + \lambda \mathbf{w}$$

Parameter update:

$$\begin{aligned} \mathbf{w}^{t+1} &= \mathbf{w}^t - \eta \frac{\partial L'}{\partial \mathbf{w}} = \mathbf{w}^t - \eta \left(\frac{\partial L}{\partial \mathbf{w}} + \lambda \mathbf{w}^t \right) \\ &= (1 - \eta \lambda) \mathbf{w}^t - \eta \frac{\partial L}{\partial \mathbf{w}} \end{aligned}$$

$(1 - \eta \lambda) \mathbf{w}^t$ always makes weight smaller (closing to zero)

Regularization

Regularization就是向你的模型加入某些规则，加入先验，缩小解空间，减小求出错误解的可能性。

$[1, 0, 0, 0]$

$l_1 = 1, l_2 = 1$

$[0.25, 0.25, 0.25, 0.25]$

$l_1 = 1, l_2 = 0.25$

L1鼓励系数稀疏

More tricks for training deep networks

- Optimizers: Momentum, AdaGrad, RMSProp, AdaDelta etc
- Learning rates
- Weight initialization
- Regularization: dropout, early stopping
- Batch normalization
- ...

实践内容

- 采用**BP**神经网络在**half moon**和**cifar10**数据集上进行分类
 1. 依照课件上的内容，实现**BP**神经网络，在**half moon**上可视化非线性分类。
 2. 在**cifar10**上，采用自己实现的**BP**神经网络来训练和测试并计算正确率。
 3. 通过调整网络每层的节点数目、**learning rate**、正则化参数、网络层数、激活函数等，来争取获得最优的分类正确率。